

Package ‘LalondeBayesBorrow’

June 11, 2025

Type Package

Title LalondeBayesBorrow: Bayesian Lalonde Framework Simulation

Version 1.1.0

Author Zhining Sui

Maintainer Zhining Sui <zsui.academics@gmail.com>

Description Implements the Bayesian Lalonde framework for simulating and evaluating go/no-go decision-making strategies in early-stage clinical trials, with a focus on leveraging historical control data. The package supports information borrowing using methods such as the Self-Adaptive Mixture (SAM) prior and power priors. It provides tools for simulating trial data for various endpoint types (continuous, binary), calculating posterior distributions, performing Bayesian inference (estimates, credible intervals, probabilities), evaluating posterior estimation metrics, and computing operating characteristics (OC) for decision rules. This package is designed to assist researchers and statisticians in the design and assessment of Bayesian adaptive trials, particularly at the Phase II/III transition.

License GPL-2

Encoding UTF-8

LazyData true

Roxygen list(markdown = TRUE)

RoxygenNote 7.3.2

Imports dplyr,
tidyr,
RBesT,
parallel,
doParallel,
foreach,
data.table,
stats,
rlang,
ggplot2,
grDevices,
glue,
scales,
stringr,
tibble

Suggests knitr,
rmarkdown,
patchwork,
kableExtra

VignetteBuilder knitr

R topics documented:

bayesian_lalonde_decision	2
calc_post_dist_metrics	4
convert_RBesT_mix	6
create_data_gen_params	7
create_pmd_summary	9
create_risk_df	9
data_gen_binary	10
data_gen_continuous	11
data_gen_DpR	13
get_all_functions	14
obtain_oc	14
plot_oc	15
plot_posterior	16
posterior_distribution	16
posterior_inference	20
posterior_prob	22
process_sim_results	23

bayesian_lalonde_decision

Bayesian Decision Framework based on Lalonde Criteria

Description

This function simulates a Bayesian trial across multiple scenarios (simulations), calculating posterior distributions for treatment and control arms (with optional historical borrowing), performing Bayesian inference (estimates, credible intervals, posterior probabilities), and applying decision rules based on the Lalonde framework.

Usage

```
bayesian_lalonde_decision(
  endpoint,
  data_summary,
  settings = NULL,
  verbose = TRUE,
  prior_params,
  lrv = 0,
  tv = 0,
  fgr = 0.2,
  fsr = 0.1,
  arm_names,
  posterior_infer = TRUE,
```

```

    Lalonde_decision = TRUE,
    EXP_TRANSFORM = FALSE
  )

```

Arguments

endpoint	A character string. The type of endpoint. Must be 'continuous' or 'binary'. This determines the likelihood and prior structure used by posterior_distribution.
data_summary	<p>A data frame. For a single analysis, this can be a single-row data frame without an nsim column. For simulations, it must contain an nsim column. Data for each arm and its parameters should be in columns prefixed by the arm name and a dot (e.g., control.n, treatment.mu_hat, control_h.count).</p> <ul style="list-style-type: none"> For 'continuous' endpoint, required columns per arm (e.g., control.) are n (sample size), mu_hat (sample mean), s (standard deviation of sample mean). For 'binary' endpoint, required columns per arm (e.g., control.) are n (sample size), count (number of responses/events). Historical arm data (e.g., control_h.n, treatment_h.count) are optional.
settings	A data frame (optional). Contains the settings for the simulation run. If provided, it will be combined with the results.
verbose	A logical value. If TRUE (default), the function prints progress messages and warnings to the console.
prior_params	<p>A named list. Contains prior distribution and borrowing parameters for each arm (treatment and control, potentially historical). Names must follow the convention [arm_name].[parameter], matching the parameters expected by posterior_distribution.</p> <ul style="list-style-type: none"> For 'continuous' arms, includes [arm_name].theta0, [arm_name].s0. For 'binary' arms, includes [arm_name].a, [arm_name].b, [arm_name].a0, [arm_name].b0. For arms with potential historical borrowing, includes [arm_name].delta, [arm_name].w, [arm_name].ess_h. Example names: treatment.delta, control.w, treatment.theta0, control.s0, treatment.a, control.b, treatment.a0, control.b0, control.ess_h.
lr_v	A scalar numeric value. The Lower Reference Value for the comparison metric (difference or ratio). Used in the Lalonde decision criteria.
tv	A scalar numeric value. The Target Value for the comparison metric. Used in the Lalonde decision criteria.
fgr	A scalar numeric value between 0 and 1. The False Go Rate. Used in the Lalonde probability and quantile-based decision criteria.
fsr	A scalar numeric value between 0 and 1. The False Stop Rate. Used in the Lalonde probability and quantile-based decision criteria.
arm_names	A character vector. A list of arm names present in the simulation. This parameter is used internally to identify which arm data columns to look for in data_summary and which prior parameters to expect in prior_params. Expected names are typically "treatment", "control", "treatment_h", "control_h".
posterior_infer	A logical value. If TRUE, compute posterior inference statistics (95% credible intervals and posterior probabilities pr_m, pr_l, pr_t). Defaults to TRUE.

Lalonde_decision

A logical value. If TRUE, apply the Lalonde decision rules based on the calculated posterior probabilities and quantiles. Defaults to TRUE. Note: Requires posterior_infer = TRUE to make decisions.

EXP_TRANSFORM

A logical value. If TRUE, results for individual arm means, credible intervals, and the comparison estimate (est_compare) are exponentiated. This implies that the underlying parameters in post_t_mix and post_c_mix (and hence mu_hat, mu in data_summary and prior_params) are on a log scale, and the comparison is interpreted as a ratio. Defaults to FALSE.

Details

It leverages posterior_distribution to calculate individual arm posteriors, posterior_inference to summarize these posteriors and their difference/ratio, and posterior_prob to calculate probabilities against thresholds.

The function is designed to work with data summarized per arm per simulation repetition, typically the output of a data generation step. Parallel processing is used to speed up calculations across simulations.

Value

A list containing three data frames, each row corresponding to a simulation repetition (nsim):

- **post_params**: Contains the raw posterior parameters (e.g., w_post, mu1_post, sigma1_post, etc.) for the treatment and control arms, prefixed with treatment. or control., and the prior weights (treatment.w_prior, control.w_prior).
- **post_est_ci**: Contains standard posterior inference (mean, SD, 95% CI) for each arm and their comparison, postfixed with _95ci. Values are exponentiated if EXP_TRANSFORM = TRUE (except SDs).
- **post_inference**: Contains inference results relevant to the Lalonde decision framework. Includes posterior probabilities (pr_m, pr_l, pr_t) and comparison quantiles derived from fgr/fsr (_lalonde). If Lalonde_decision = TRUE and posterior_infer = TRUE, includes decision columns (decision_pr, decision_ci).

See Also

posterior_distribution, posterior_inference, posterior_prob, convert_RBesT_mix, RBesT::mix, RBesT::mixnorm, RBesT::mixbeta, RBesT::pmixdiff, RBesT::qmixdiff

calc_post_dist_metrics

Calculate Posterior Distribution Metrics

Description

This function calculates various evaluation metrics for the posterior distribution, including the average bias, standard error, empirical standard deviation, and nominal coverage probability. For the "g-score" endpoint, it calculates the log-median ratio, and for the "OR" endpoint, it calculates the rate difference.

Usage

```
calc_post_dist_metrics(endpoint, true_value, post_est_ci, remove.na = FALSE)
```

Arguments

endpoint	A string. Specifies the type of endpoint ("g-score" or "OR" or "continuous").
true_value	A scalar. The true median ratio for the "g-score" endpoint, the true rate difference for the "OR" endpoint, or the true mean difference for the "continuous" endpoint.
post_est_ci	A data frame. Contains posterior inference obtained from <code>bayesian_lalonde_decision()</code> , including posterior estimates, standard errors, and 95% credible intervals for each simulation repetition.
remove.na	Logical. Whether to remove NA values before calculation. Defaults to FALSE.

Details

For the "g-score" endpoint, it calculates the log-median ratio, and for the "OR" endpoint, it calculates the rate difference.

- Average bias
- Standard error
- Empirical standard deviation
- Nominal coverage probability

Value

A data frame containing the evaluation metrics for the posterior distribution, including:

- `bias_avg`: The average bias of the estimate (specific column names like `bias_avg_median_ratio1`, `bias_avg_median_ratio2`, `bias_avg_rate_diff`, `bias_avg_mean_diff` depend on the endpoint).
- `sd_avg`: The average standard error of the estimate (`sd_compare`).
- `sd_empirical`: The empirical standard deviation of the estimate (`est_compare`).
- `cp`: The nominal coverage probability (95%).
- Additional metrics specific to the "g-score" endpoint (see `eval_gscore_approx_dist()`).

See Also

`bayesian_lalonde_decision`, `eval_gscore_approx_dist`

convert_RBesT_mix	<i>Convert to RBesT Mixture Distribution</i>
-------------------	--

Description

This function generates a mixture distribution object compatible with the RBesT package from a data frame containing component parameters. It supports the creation of both normal (mixnorm) and beta (mixbeta) mixture distributions based on the specified endpoint type. The function is designed to handle mixture models with up to two components based on the provided structure of the input data frame.

Usage

```
convert_RBesT_mix(post, endpoint)
```

Arguments

post	<p>A data frame or a list of parameters. This object should contain the parameters defining the mixture components. The expected structure depends on the endpoint type:</p> <ul style="list-style-type: none"> For endpoint = "continuous" (Normal mixture): Expected columns/elements are w (weight of the first component), mu1 (mean of the first component), sigma1 (standard deviation of the first component), mu2 (mean of the second component), and sigma2 (standard deviation of the second component). The weight of the second component is implicitly 1 - w. For endpoint = "binary" (Beta mixture): Expected columns/elements are w (weight of the first component), a1 (alpha parameter of the first component), b1 (beta parameter of the first component), a2 (alpha parameter of the second component), and b2 (beta parameter of the second component). The weight of the second component is implicitly 1 - w. The function currently assumes a single row/set of parameters for the mixture, representing either a one- or two-component mixture.
endpoint	<p>A character string. Specifies the type of endpoint and thus the type of mixture distribution to create. Must be either "continuous" for a normal mixture (mixnorm) or "binary" for a beta mixture (mixbeta).</p>

Value

An RBesT mixture distribution object, specifically of class mixnorm if endpoint is "continuous" or mixbeta if endpoint is "binary". Returns an error if the endpoint is not specified correctly or if required parameters are missing or invalid.

See Also

[mixnorm](#), @seealso [mixbeta](#), @seealso [mix](#)

Examples

```
# Example for a continuous endpoint (Normal mixture)
post_continuous <- data.frame(w = 0.6, mu1 = 10, sigma1 = 2, mu2 = 15, sigma2 = 3)
mix_norm <- convert_RBesT_mix(post_continuous, endpoint = "continuous")
```

```

print(mix_norm)

# Example for a binary endpoint (Beta mixture)
post_binary <- list(w = 0.8, a1 = 5, b1 = 15, a2 = 10, b2 = 10)
mix_beta <- convert_RBES_T_mix(post_binary, endpoint = "binary")
print(mix_beta)

# Example for a single-component mixture (weight = 1)
post_single <- data.frame(w = 1, mu1 = 12, sigma1 = 1.5, mu2 = NA, sigma2 = NA)
mix_single <- convert_RBES_T_mix(post_single, endpoint = "continuous")
print(mix_single)

# Example for a single-component mixture (weight = 0)
post_single_comp2 <- list(w = 0, mu1 = NA, sigma1 = NA, mu2 = 8, sigma2 = 2.5)
mix_single_comp2 <- convert_RBES_T_mix(post_single_comp2, endpoint = "continuous")
print(mix_single_comp2)

```

create_data_gen_params

Create Data Generation Parameters

Description

This function processes and validates a set of parameters for generating data for different study arms (treatment, control, and their historical counterparts). It checks for missing parameters, replaces them with NA while issuing warnings, and structures the valid parameters into a nested list format suitable for data simulation functions. The function supports various endpoint types, each requiring a specific set of parameters and potentially applying transformations (like logarithm for means in "g-score").

Usage

```
create_data_gen_params(params, endpoint)
```

Arguments

params	A list. Contains the parameters for data generation. Expected parameters follow a naming convention: [arm_prefix]_[parameter_name], where [arm_prefix] can be trt (treatment), ctrl (control), trt_h (historical treatment), or ctrl_h (historical control). Required [parameter_name]s depend on the endpoint. Examples include n (sample size), p (probability), mu (mean), sigma (standard deviation).
endpoint	A character string. Specifies the type of endpoint. Must be one of "continuous", "binary", "DpR", or "g-score".

Details

Arms (treatment, control, historical treatment, historical control) are included in the output list only if all their required parameters for the specified endpoint are present (not NULL and not resulting in NA after initial checks), and if the sample size n is positive.

The interpretation of parameters depends on the endpoint:

- **continuous:** Parameters for a normal continuous outcome include a mean (μ) and a standard deviation (σ). Required parameters per arm: n , μ , σ .
- **binary:** Parameters for a binary outcome include a probability of the event occurring (p). Required parameters per arm: n , p .
- **DpR** (Depth of Response): Parameters include a probability of complete tumor shrinkage (p), a mean (μ), and a standard deviation (σ) for the continuous outcome part. Required parameters per arm: n , p , μ , σ .
- **g-score:** Parameters for zero-inflated continuous g-scores include a probability of observing a zero value (p), a mean (μ , log-transformed internally) and a standard deviation (σ) for the non-zero part. Required parameters per arm: n , p , μ , σ .

Value

A named list where each element corresponds to a study arm (treatment, control, treatment_h, control_h). Each arm's element is a list containing its data generating parameters (n , p , μ , σ , etc., depending on the endpoint) and its name. Arms for which required parameters were missing (NULL or resulted in NA) or n was non-positive are excluded from the list. Missing parameters for included arms are reported via warnings and set to NA.

Examples

```
# Example for 'continuous' endpoint with missing historical data
params_continuous <- list(
  trt_n = 150, trt_mu = 75, trt_sigma = 10,
  ctrl_n = 150, ctrl_mu = 70, ctrl_sigma = 11,
  trt_h_n = 80, trt_h_mu = 76 # Missing sigma, this arm will be excluded
)
data_params_continuous <- create_data_gen_params(params_continuous, "continuous")
str(data_params_continuous)

# Example for 'binary' endpoint
params_binary <- list(
  trt_n = 200, trt_p = 0.7,
  ctrl_n = 200, ctrl_p = 0.5,
  ctrl_h_n = 100, ctrl_h_p = 0.55
)
data_params_binary <- create_data_gen_params(params_binary, "binary")
str(data_params_binary)

# Example for 'DpR' endpoint
params_DpR <- list(
  trt_n = 120, trt_p = 0.9, trt_mu = 5, trt_sigma = 1.2, # p=prob of complete shrinkage
  ctrl_n = 120, ctrl_p = 0.8, ctrl_mu = 4.5, ctrl_sigma = 1.5
)
data_params_DpR <- create_data_gen_params(params_DpR, "DpR")
str(data_params_DpR)

# Example for 'g-score' endpoint
params_gscore <- list(
  trt_n = 100, trt_p = 0.2, trt_mu = 1.5, trt_sigma = 0.5, # p=prob of zero g-score
  ctrl_n = 100, ctrl_p = 0.4, ctrl_mu = 1.2, ctrl_sigma = 0.6
)
data_params_gscore <- create_data_gen_params(params_gscore, "g-score")
str(data_params_gscore)
```

create_pmd_summary	<i>Create a Data Frame of Posterior Mean Differences (PMD)</i>
--------------------	--

Description

This function calculates the Posterior Mean Difference (PMD) by comparing any two specified borrowing scenarios from a simulation run. It correctly handles both one-to-one comparisons (e.g., 'SAM Prior' vs. 'Fixed Weight') and one-to-many comparisons against a common baseline (e.g., comparing multiple borrowing scenarios to a single 'No Borrowing' run).

Usage

```
create_pmd_summary(
  post_inference_all,
  borrow_to_compare = c("Yes", "No"),
  pmd_var = "est2_lalonde",
  group_vars = NULL
)
```

Arguments

post_inference_all	A data frame containing the posterior inference results from a simulation run, typically from <code>process_sim_results()</code> .
borrow_to_compare	A character vector of length two specifying the two borrowing conditions from the 'borrow' column to compare (e.g., <code>c("Yes", "No")</code>). The difference will be calculated as <code>borrow_to_compare[1] - borrow_to_compare[2]</code> .
pmd_var	The name of the single posterior estimate variable to use for calculating the PMD. Defaults to "est2_lalonde".
group_vars	A character vector of variable names to group by. If NULL, the function will automatically use all setting columns.

Value

A data frame summarizing the mean and standard deviation of the PMD for each relevant combination of settings.

create_risk_df	<i>Create a Data Frame of Operating Characteristic Risks</i>
----------------	--

Description

This function takes the operating characteristics from a simulation run and creates a data frame summarizing the False Go Rate (FGR) and False Stop Rate (FSR) for each simulation scenario.

Usage

```
create_risk_df(oc_all, lrv, tv, group_vars = NULL)
```

Arguments

oc_all	A data frame from process_sim_results().
lr_v	The Lower Reference Value used in the simulation.
tv	The Target Value used in the simulation.
group_vars	A character vector of variable names to group by. If NULL (default), the function will automatically use all setting columns.

Value

A data frame with the FGR and FSR for each simulation scenario.

data_gen_binary	<i>Generate Binary Response Data</i>
-----------------	--------------------------------------

Description

This function simulates binary endpoint data (0/1) for multiple study arms across several simulation repetitions. Each arm's data is generated based on the specified probabilities of response using a binomial distribution. The function returns individual-level data, true response rates, and response frequencies per arm per simulation.

Usage

```
data_gen_binary(
  n_arms = 2,
  nsim = 10,
  n_list = NULL,
  prob_list = NULL,
  arm_names = NULL,
  seed = NULL
)
```

Arguments

n_arms	Number of arms (optional, inferred from arm_names if not provided).
nsim	A scalar. The number of repetitions of the simulation to be performed.
n_list	A named list. Each element represents the number of patients in the corresponding arm (names must match arm_names). Must contain positive integers.
prob_list	A named list. Each element represents the probability of response (binary outcome = 1) for the corresponding arm (names must match arm_names). Must contain numeric values between 0 and 1.
arm_names	A character vector of arm names.
seed	An integer or NULL. If an integer, a seed is set for reproducibility of the random number generation across parallel workers. Defaults to NULL, which means no explicit seed is set beyond R's default behavior.

Value

A list containing:

- `data`: A `data.table` with columns `nsim`, `id`, `arm`, `resp`, containing all simulated binary outcomes (0/1) for all arms and all repetitions.
- `true_value`: A data frame of true response rates for each arm, along with the difference in response rates between the treatment and control arms (if both are present).
- `summary`: A `data.table` of response frequencies per arm per simulation, # Corrected from 'freq' to 'summary' based on code with columns for `nsim`, `arm`, `count` (number of responses), and `n` (total patients).

See Also

`bayesian_lalonde_decision`

<code>data_gen_continuous</code>	<i>Generate continuous outcome data for multiple arms across multiple simulations</i>
----------------------------------	---

Description

Generates continuous outcome data by sampling from Normal distributions for each arm in each simulation replicate. It produces both the individual-level data and arm-specific summary statistics (sample size, mean, standard deviation, standard error). The function supports parallel processing for faster execution across simulations and includes options for reproducibility via setting a random seed.

Usage

```
data_gen_continuous(
  nsim = 10,
  n_list,
  mu_list,
  sigma_list,
  arm_names,
  seed = NULL
)
```

Arguments

<code>nsim</code>	A single positive integer. Number of simulations to run.
<code>n_list</code>	A named list of positive integers. Sample sizes for each arm. Names must exactly match <code>arm_names</code> .
<code>mu_list</code>	A named list of numeric values. True means for the Normal distribution for each arm. Names must exactly match <code>arm_names</code> .
<code>sigma_list</code>	A named list of non-negative numeric values. True standard deviations for the Normal distribution for each arm. Names must exactly match <code>arm_names</code> .
<code>arm_names</code>	A character vector. Names of the arms in the simulation. Used for naming outputs and matching list elements.
<code>seed</code>	An integer or <code>NULL</code> . If an integer, a seed is set for reproducibility of the random number generation across parallel workers. Defaults to <code>NULL</code> , which means no explicit seed is set beyond R's default behavior.

Value

A list containing:

- `data`: A `data.table` with columns `nsim` (simulation replicate ID), `id` (patient ID within the replicate and arm), `arm` (arm name), `value` (simulated continuous outcome). Contains all individual simulated values.
- `summary`: A `data.table` with columns `nsim` and arm-specific summary statistics calculated from the simulated data for each arm and simulation replicate. Columns are in wide format, named `[arm_name].n`, `[arm_name].mu_hat` (sample mean), `[arm_name].sd` (sample standard deviation), and `[arm_name].s` (sample standard error of the mean, sd/\sqrt{n}).
- `true_value`: A `data.frame` with the true input mean values per arm (columns `[arm_name].mean_true`) and a true mean difference between 'treatment' and 'control' arms if both exist (column `compare_true`).

See Also

[bayesian_lalonde_decision](#)

Examples

```
# Example 1: Generate data for 2 arms over 5 simulations
n_list <- list(treatment = 50, control = 60)
mu_list <- list(treatment = 10, control = 8)
sigma_list <- list(treatment = 3, control = 3.5)
arm_names <- c("treatment", "control")

sim_data <- data_gen_continuous(nsim = 5, n_list, mu_list, sigma_list, arm_names)

# Display structure of results
str(sim_data)

# Display first few rows of individual data
print(head(sim_data$data))

# Display the summary data
print(sim_data$summary)

# Display the true values
print(sim_data$true_value)

# Example 2: Generate data with a seed for reproducibility
sim_data_reproducible <- data_gen_continuous(nsim = 3, n_list, mu_list,
                                             sigma_list, arm_names, seed = 123)
print(sim_data_reproducible$summary) # Note the summary values

sim_data_reproducible_again <- data_gen_continuous(nsim = 3, n_list, mu_list,
                                                  sigma_list, arm_names, seed = 123)
print(sim_data_reproducible_again$summary) # Values should match the previous run
```

data_gen_DpR	<i>Generate Constrained Depth of Response (DpR) data for multiple arms across multiple simulations</i>
--------------	--

Description

Generate Constrained Depth of Response (DpR) data for multiple arms across multiple simulations

Usage

```
data_gen_DpR(n_arms, nsim, n_list, p_list, arm_names, mu_list, sigma_list)
```

Arguments

n_arms	Number of arms (can be inferred from arm_names, but kept for consistency with template)
nsim	Number of simulations
n_list	Named list of sample sizes per arm (names should match arm_names)
p_list	Named list of proportion of -1 per arm (names should match arm_names)
arm_names	Character vector of arm names
mu_list	Named list of means per arm (theta_x) (names should match arm_names)
sigma_list	Named list of standard deviations per arm (sigma_x) (names should match arm_names)

Value

A list containing:

- data: A data.table with columns nsim, id, arm, value, containing all simulated DpR values.
- summary: A data.table with columns nsim and arm-specific summary statistics (n, mu_hat, sd, s).
- true_value: A data frame with true mean differences (currently only 'treatment' vs 'control').

See Also

bayesian_lalonde_decision

get_all_functions	<i>Retrieve All Functions from an Environment</i>
-------------------	---

Description

This function returns the names of all functions within a specified environment.

Usage

```
get_all_functions(env = environment())
```

Arguments

env	An environment. The environment to search for functions. The default is the current environment.
-----	--

Value

A character vector. Names of all functions found in the specified environment.

obtain_oc	<i>Obtain Operating Characteristics (OC)</i>
-----------	--

Description

This function calculates the operating characteristics (OC) by computing the proportion of decisions made based on the posterior probability and the credible interval from a set of decision results. It combines these proportions with the provided settings data to return a comprehensive summary.

Usage

```
obtain_oc(decisions)
```

Arguments

decisions	A data frame. Contains the decisions made in each simulation or trial, with columns for posterior probability-based decisions (decision_pr) and credible interval-based decisions (decision_ci).
-----------	--

Value

A data frame that includes the original settings along with the calculated proportions of decisions based on the posterior probability and credible interval. The output includes:

- proportion_pr: Proportion of decisions made based on the posterior probability.
- proportion_ci: Proportion of decisions made based on the credible interval.

plot_oc

*Plot Operating Characteristics (OC)***Description**

Creates a stacked plot of the operating characteristics, showing the probabilities of "Go", "No-Go", and "Consider" decisions for each simulation scenario.

Usage

```
plot_oc(oc_data, x_var, facet_formula = NULL, plot_type = "stepwise", ...)
```

Arguments

oc_data	A data frame containing the operating characteristics, typically from <code>process_sim_results()</code> . Must contain columns for the decision proportions (e.g., <code>proportion_pr</code>) and the decision types (e.g., <code>decision_pr</code>).
x_var	A character string specifying the variable for the x-axis (e.g., <code>"control.n"</code> or <code>"control.p.diff"</code>).
facet_formula	A formula for faceting the plot, to split it into panels based on other setting variables (e.g., <code>~ borrowing</code> or <code>true_value.compare_true ~ borrowing</code>).
plot_type	A character string specifying the type of plot. Either <code>"stepwise"</code> for a bar plot with text or <code>"smooth"</code> for a stacked area plot. Defaults to <code>"stepwise"</code> .
...	Additional arguments passed to <code>ggplot2</code> functions for further customization (e.g., <code>labs()</code> , <code>theme()</code>).

Value

A `ggplot` object.

Examples

```
# Assuming 'processed_results' is the output from process_sim_results()
# and contains an 'oc_all' data frame.

# plot_oc(
#   oc_data = processed_results$oc_all,
#   x_var = "control.n",
#   facet_formula = true_value.compare_true ~ borrowing,
#   labs(
#     title = "Operating Characteristics by Sample Size and Borrowing",
#     x = "Sample Size per Arm",
#     y = "Probability of Decision"
#   )
# )
```

plot_posterior	<i>Plot Posterior Distributions for a Single Run</i>
----------------	--

Description

This function visualizes and compares the posterior distributions for the treatment arm, the control arm (without borrowing), and the hybrid control arm (with borrowing) from a single simulation run.

Usage

```
plot_posterior(
  post_params,
  post_inference,
  endpoint,
  title_format = NULL,
  x_range = NULL,
  ...
)
```

Arguments

post_params	A data frame with exactly two rows: one for a "borrowing" scenario and one for a "no borrowing" scenario from the same simulation run.
post_inference	A data frame with the corresponding posterior inference results for the "borrowing" scenario. This should be filtered to a single row.
endpoint	A character string specifying the endpoint type: "binary" or "continuous".
title_format	A character string that defines the format for the plot title. You can use {column_name} syntax to insert values from the post_params data frame (e.g., "Control: {control.count}/{control.count}"). If NULL, a default title will be used.
x_range	A numeric vector of length two specifying the range for the x-axis. If NULL, a reasonable default will be used.
...	Additional arguments passed to ggplot2 functions for further customization (e.g., labs(), theme()).

Value

A ggplot object visualizing the posterior distributions.

posterior_distribution	<i>Obtain Posterior Distribution for Single Arm</i>
------------------------	---

Description

This function calculates the posterior distribution for one-arm trial data with either continuous or binary endpoints, incorporating current data and optionally borrowing from historical data. It supports the SAM prior (Self-Adaptive Mixture) for automatic conflict detection when borrowing is enabled ($w = \text{NULL}$) and allows specification of a fixed mixture weight for the informative component (w is numeric). A power prior mechanism is included via the `ess_h` parameter to control the influence of historical data. Additionally, a gating mechanism is applied: if the prior-data conflict exceeds a threshold (`delta_gate`), borrowing is disabled regardless of the SAM weight.

Usage

```
posterior_distribution(
  endpoint,
  current,
  historical = NULL,
  delta_SAM = NULL,
  delta_gate = NULL,
  w = NULL,
  a = 0.01,
  b = 0.01,
  a0 = 0.01,
  b0 = 0.01,
  theta0 = 0,
  s0 = 100,
  ess_h = NULL
)
```

Arguments

<code>endpoint</code>	A string. The type of endpoint, either 'continuous' (assumes Normal likelihood and Normal-Inverse-Gamma conjugacy on mean/variance) or 'binary' (assumes Binomial likelihood and Beta conjugacy).
<code>current</code>	A named list or vector containing current trial data: <ul style="list-style-type: none"> For 'continuous': must include <code>n</code> (number of patients), <code>mu_hat</code> (sample mean), <code>s</code> (sample standard deviation the mean). For 'binary': must include <code>n</code> (number of patients), <code>count</code> (number of responses/events). Sample size <code>n</code> must be positive.
<code>historical</code>	A named list or vector containing historical trial data in the same format as the current parameter. If <code>NULL</code> (default), no historical data is used, and borrowing is disabled (w is effectively 0). If provided, historical sample size <code>n</code> must be non-negative.
<code>delta_SAM</code>	A scalar positive value, or <code>NULL</code> . The clinical significant difference threshold used in the SAM prior calculation to detect prior-data conflict. Only used if w is <code>NULL</code> . The interpretation depends on the endpoint: for 'continuous', it's on the mean scale; for 'binary', it's on the probability scale. If <code>NULL</code> and w is <code>NULL</code> , default values (0.2 for continuous, 0.1 for binary) are used.
<code>delta_gate</code>	A scalar positive value or <code>NULL</code> . Threshold for the gating mechanism: if prior-data conflict exceeds this, borrowing is disabled. If <code>NULL</code> , gate is not applied.
<code>w</code>	A scalar value between 0 and 1, or <code>NULL</code> .

	<ul style="list-style-type: none"> • If a numeric value (0 to 1): Specifies a fixed initial weight for the informative (historical) component of the mixture prior. Borrowing occurs according to this fixed weight. • If NULL (default): The SAM prior is used. The initial weight for the informative component is automatically calculated based on the agreement between historical and current data, using <code>delta</code>. Note: If historical is NULL or <code>ess_h</code> is 0, <code>w</code> is effectively forced to 0, and no borrowing occurs regardless of the input <code>w</code>.
<code>a</code>	A scalar positive value. The alpha parameter for the <i>non-informative</i> Beta prior component (used only for endpoint = "binary"). Defaults to 0.01.
<code>b</code>	A scalar positive value. The beta parameter for the <i>non-informative</i> Beta prior component (used only for endpoint = "binary"). Defaults to 0.01.
<code>a0</code>	A scalar positive value. The base alpha parameter for the <i>informative</i> Beta prior component (used only for endpoint = "binary"). This parameter (along with <code>b0</code>) serves as a base prior which is updated by the discounted historical data to form the informative prior component parameters. Defaults to 0.01.
<code>b0</code>	A scalar positive value. The base beta parameter for the <i>informative</i> Beta prior component (used only for endpoint = "binary"). Defaults to 0.01.
<code>theta0</code>	A scalar. The mean parameter for the <i>non-informative</i> Normal prior component (used only for endpoint = "continuous"). Defaults to 0.
<code>s0</code>	A scalar positive value. The standard deviation parameter for the <i>non-informative</i> Normal prior component (used only for endpoint = "continuous"). Defaults to 100.
<code>ess_h</code>	A scalar non-negative finite value, or NULL. The Effective Sample Size to use for the historical data. This implements a power prior approach, where the historical data likelihood is raised to the power of <code>ess_h / nh</code> . If NULL (default), the power is 1 (full historical data influence). If 0, historical data is fully discounted, equivalent to no borrowing.

Value

A list containing the results of the posterior calculation:

- `w_prior`: The initial mixture weight used for the informative prior component *before* updating with current data. This is either the input `w` or the weight calculated by the SAM prior if `w` was NULL. Returns 0 if no historical data was provided, historical `n` was 0, or `ess_h` was 0.
- `post`: A named list containing the parameters of the resulting two-component mixture posterior distribution, suitable for use with `RBesT`:
 - `w`: The *posterior* weight of the informative component. This is calculated by updating the `w_prior` with the current data via Likelihood Ratio Test.
 - For continuous (mixnorm format): `mu1`, `sigma1` (mean and SD of the informative posterior component), `mu2`, `sigma2` (mean and SD of the non-informative posterior component). These will be `NA_real_` if the corresponding posterior component has zero weight (`post$w` is 0 or 1).
 - For binary (mixbeta format): `a1`, `b1` (parameters of the informative Beta posterior component), `a2`, `b2` (parameters of the non-informative Beta posterior component). These will be `NA_real_` if the corresponding posterior component has zero weight (`post$w` is 0 or 1). Returns an error if inputs are invalid or required parameters are missing.

posterior_inference, posterior_prob, mix, mixnorm, mixbeta

[illegible]

```
str(post_noborrow_bin)
```

posterior_inference *Perform Posterior Inference*

Description

This function calculates key summary statistics (posterior mean, standard deviation, and credible intervals) for one or two RBesT mixture posterior distributions. When two distributions are provided (e.g., for treatment and control arms), it also computes these statistics for the difference or ratio between them. The function supports exponentiating the results, useful when the original parameter of interest is on a log scale (e.g., log-odds, log-mean ratio).

Usage

```
posterior_inference(post1, post2 = NULL, quantiles, EXP_TRANSFORM = FALSE)
```

Arguments

post1	An RBesT mixture distribution object (class mix, mixnorm, mixbeta, etc.). Represents the posterior distribution for the first entity (e.g., treatment arm).
post2	An RBesT mixture distribution object (class mix, mixnorm, mixbeta, etc.), or NULL. Represents the posterior distribution for the second entity (e.g., control arm). If NULL (default), inference is performed only for post1.
quantiles	A numeric vector of length 2 specifying the lower and upper quantiles for the credible interval (e.g., c(0.025, 0.975) for a 95% credible interval). Must be sorted and within (0, 1).
EXP_TRANSFORM	A logical value. If TRUE, the calculated mean, credible interval bounds, and the comparison estimate are exponentiated ($\exp(x)$). This is typically used when the mixture distribution models a parameter on the log scale (e.g., log-mean, log-odds). Note that the standard deviation (sd1, sd2, sd_compare) is always calculated on the original scale of the input mixture distributions.

Value

A data frame with exactly one row. Contains the following columns:

- est1: Posterior mean for post1. Exponentiated if EXP_TRANSFORM = TRUE.
- sd1: Posterior standard deviation for post1. Always on the original scale of post1.
- ci1_l: Lower bound of the credible interval for post1. Exponentiated if EXP_TRANSFORM = TRUE.
- ci1_u: Upper bound of the credible interval for post1. Exponentiated if EXP_TRANSFORM = TRUE.
- est2: Posterior mean for post2 (if post2 is not NULL). Exponentiated if EXP_TRANSFORM = TRUE. NA if post2 is NULL.
- sd2: Posterior standard deviation for post2 (if post2 is not NULL). Always on the original scale of post2. NA if post2 is NULL.

- `ci2_l`: Lower bound of the credible interval for `post2` (if `post2` is not NULL). Exponentiated if `EXP_TRANSFORM = TRUE`. NA if `post2` is NULL.
- `ci2_u`: Upper bound of the credible interval for `post2` (if `post2` is not NULL). Exponentiated if `EXP_TRANSFORM = TRUE`. NA if `post2` is NULL.
- `est_compare`: Estimated treatment effect. If `EXP_TRANSFORM = FALSE`, this is `est1 - est2` (difference of means/parameters). If `EXP_TRANSFORM = TRUE`, this is `exp(est1) / exp(est2)` (ratio of exponentiated means/parameters). NA if `post2` is NULL.
- `sd_compare`: Standard deviation of the *difference* (`post1 - post2`) on the original scale. Calculated as $\sqrt{sd1^2 + sd2^2}$, assuming independence. This is *not* the standard deviation of the ratio if `EXP_TRANSFORM = TRUE`. NA if `post2` is NULL.
- `compare_ci_l`: Lower bound of the credible interval for the treatment effect (`post1 - post2`). Exponentiated if `EXP_TRANSFORM = TRUE`. NA if `post2` is NULL.
- `compare_ci_u`: Upper bound of the credible interval for the treatment effect (`post1 - post2`). Exponentiated if `EXP_TRANSFORM = TRUE`. NA if `post2` is NULL.

See Also

[bayesian_lalonde_decision](#), [posterior_distribution](#), [posterior_prob](#), [qmix](#), [qmixdiff](#)

Examples

```
# Assuming RBest is installed and loaded
library(RBest) # Load RBest for example
# Example 1: Single arm, continuous (Normal mixture)
post_single_norm <- mixnorm(c(0.6, 10, 2), c(0.4, 15, 3))
inference_single <- posterior_inference(post1 = post_single_norm, quantiles = c(0.025, 0.975))
print(inference_single)

# Example 2: Two arms, continuous (Difference)
post_ctrl_norm <- mixnorm(c(1, 12, 2.5)) # Control arm posterior
inference_diff <- posterior_inference(post1 = post_single_norm, post2 = post_ctrl_norm,
                                     quantiles = c(0.025, 0.975))
print(inference_diff)

# Example 3: Two arms, binary (Beta mixture), inference on probability scale
post_trt_beta <- mixbeta(c(0.7, 10, 5), c(0.3, 2, 8)) # Treatment posterior for probability
post_ctrl_beta <- mixbeta(c(1, 8, 12)) # Control posterior for probability
inference_prob_diff <- posterior_inference(post1 = post_trt_beta, post2 = post_ctrl_beta,
                                          quantiles = c(0.05, 0.95))
print(inference_prob_diff)

# Example 4: Two arms, log-transformed data, inference on ratio scale
# Assume post_log_trt and post_log_ctrl are posteriors on the log scale
post_log_trt <- mixnorm(c(1, log(10), 0.5))
post_log_ctrl <- mixnorm(c(1, log(5), 0.6))
inference_log_ratio <- posterior_inference(post1 = post_log_trt, post2 = post_log_ctrl,
                                          quantiles = c(0.025, 0.975), EXP_TRANSFORM = TRUE)

# Note: est_compare is exp(log(10) - log(5)) = 10/5 = 2
# credible interval is exp(CI_log_diff) = exp(CI_log_trt - CI_log_ctrl) = CI_ratio
print(inference_log_ratio)
```

posterior_prob

*Calculate Posterior Probability***Description**

This function calculates the posterior probability of a parameter estimate (or the difference between two parameter estimates) falling within a specified range, based on provided RBesT mixture posterior distributions. For a single posterior distribution, it computes the probability of the parameter being greater than, less than, or between specified values. For two posterior distributions, it computes the probability of the difference between the first and second distributions falling within the range.

Usage

```
posterior_prob(
  range_type = "greater",
  post1,
  post2 = NULL,
  value,
  value2 = NULL
)
```

Arguments

range_type	A character string. Specifies the type of comparison range. Must be one of 'greater', 'less', or 'between'. Defaults to 'greater'.
post1	An RBesT mixture distribution object (class mix, mixnorm, mixbeta, etc.). Represents the first posterior distribution.
post2	An RBesT mixture distribution object (class mix, mixnorm, mixbeta, etc.). Optional second posterior distribution object for pairwise comparison of post1 - post2. If NULL (default), the probability is calculated for post1 itself.
value	A scalar numeric value. The threshold used for comparison. If range_type is "greater" or "less", this is the single threshold. If range_type is "between", this is the lower bound of the interval.
value2	A scalar numeric value. The upper bound for the "between" comparison. This parameter is required if range_type is "between".

Value

A scalar numeric value representing the calculated posterior probability, between 0 and 1. Returns NA if there is an error during the probability calculation (e.g., pmixdiff fails for two distributions).

See Also

[posterior_distribution](#), [posterior_inference](#), [pmix](#), [pmixdiff](#)

Examples

```
# Assuming RBest is installed and loaded
library(RBest) # Load RBest for example
# Example 1: Single arm - Probability of being greater than a value
post_single_norm <- mixnorm(c(0.6, 10, 2), c(0.4, 15, 3))
prob_greater <- posterior_prob(post1 = post_single_norm, value = 12, range_type = "greater")
cat("P(post_single_norm > 12):", prob_greater, "\n")

# Example 2: Single arm - Probability of being less than a value
prob_less <- posterior_prob(post1 = post_single_norm, value = 11, range_type = "less")
cat("P(post_single_norm < 11):", prob_less, "\n")

# Example 3: Single arm - Probability of being between two values
prob_between <- posterior_prob(post1 = post_single_norm, value = 11,
                              range_type = "between", value2 = 14)
cat("P(11 < post_single_norm < 14):", prob_between, "\n")

# Example 4: Two arms - Probability of difference > 0 (post1 > post2)
post_ctrl_norm <- mixnorm(c(1, 12, 2.5)) # Control arm posterior
prob_diff_greater_0 <- posterior_prob(post1 = post_single_norm, post2 = post_ctrl_norm,
                                     value = 0, range_type = "greater")
cat("P(post_single_norm - post_ctrl_norm > 0):", prob_diff_greater_0, "\n")

# Example 5: Two arms - Probability of difference < a value
prob_diff_less_neg1 <- posterior_prob(post1 = post_single_norm, post2 = post_ctrl_norm,
                                     value = -1, range_type = "less")
cat("P(post_single_norm - post_ctrl_norm < -1):", prob_diff_less_neg1, "\n")

# Example 6: Two arms - Probability of difference between two values
prob_diff_between <- posterior_prob(post1 = post_single_norm, post2 = post_ctrl_norm,
                                   value = -2, range_type = "between", value2 = 1)
cat("P(-2 < post_single_norm - post_ctrl_norm < 1):", prob_diff_between, "\n")
```

process_sim_results *Process and Consolidate Simulation Results*

Description

This function takes a list of results from multiple simulation scenarios (each being the output of `bayesian_lalonde_decision`) and consolidates them into a final, tidy list of data frames ready for analysis and visualization.

Usage

```
process_sim_results(bayes_results)
```

Arguments

`bayes_results` A list where each element is a list of results from one call to `bayesian_lalonde_decision`.

Value

A list of consolidated data frames: `post_params_all`, `post_est_ci_all`, `post_inference_all`, and `oc_all`.